

Fruit classification with a CNN

This notebook will serve as implementation of the API that you have created in your "Code" folder. You will write functions in the ".py" files and use them here.

We will be using "Fruits" dataset present in PyTorch and train a convolutional neural network (CNN) to classify digits.

What is expected from this

This notebook should be used to present your work. You should explain wherever necessary (but also not too much) about what you did and why you did it. You should explain things like hyper parameter settings (even if it was provided before hand to you by us), training performance and testing performance of the model. You should reason why your model is working fine and not overfitting.

Since numbers don't are an argot, you should also use visualizations wherever possible. You can visualize things like **loss curve**, show **confusion matrix** and since this is a CNN you can also consider **advance techniques like gradcam**, etc.

You can also use techniques that allow for faster training, assuage problems involving vanishing and exploding gradients.

Finally, you can show some manual verifications by displaying and making predictions on random test examples.

Absolutely required items?

1. First of all, import the libraries and the dataset. Divide the dataset into test and train.
2. Next, show dataset samples and distribution of different type of data. For example, in case of "Fruits Dataset" you can show some random images and their labels. Also, show distribution of each class of images.
3. Next, perform required transformations (also **data augmentation**) on "Fruits dataset" (normalization, resizing, grayscaling, if required, etc.) using torchvision transforms.
4. Create required dataloaders with PyTorch and use the module dataset we created to load data in mini-batches.
5. Train the model, show loss and accuracy at each step of operation.
6. Plot the **loss curve for both train and validation phase**.
7. Pick some manual random images (probably 7-10) from test dataset and predict their values showing **expected and actual result**.

```
In [ ]: #Imports
from utils import dataset
import torch.nn as nn
from utils import dataset
import torchvision.transforms as transforms
import torch
import numpy as np
from torch.autograd import Variable
from PIL import Image
import numpy as np
import os
import pandas as pd
from model import FNet
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
# import other required libraries here
```

```
In [ ]: # loading dataframes using dataset module
df, df_train, df_test = dataset.create_and_load_meta_csv_df(dataset_path=
```

```
In [ ]: s1 = np.floor(0.8 * len(df_train)) #conversion as index won't take a float
train_data = df_train.iloc[0:int(s1)]
validation_data = df_train.iloc[int(s1):]

#Resetting index so that whole dataframe can be used for iteration
train_data = train_data.reset_index(drop=True)
validation_data = validation_data.reset_index(drop=True)
```

```
In [ ]: # using dataframes, pytorch and torchvision to transform data. Also, use
#Transformation:
# ToPILImage :: will convert image to pil image while preserving range
# ToTensor :: will convert every image pixel float number between range

data_transforms = {
    'train': transforms.Compose([
        transforms.ToPILImage(), transforms.ToTensor()
    ]),
    'val': transforms.Compose([
        transforms.ToPILImage(), transforms.ToTensor()
    ]),
    'test': transforms.Compose([
        transforms.ToPILImage(), transforms.ToTensor()
    ])
}

image_datasets = {'train': dataset.ImageDataset(df_train, transform=data_train),
                  'val': dataset.ImageDataset(df_test, transform=data_train),
                  'test': dataset.ImageDataset(df_test, transform = data_train)}

train_dataset = image_datasets['train']
test_dataset = image_datasets['test']
validation_dataset = image_datasets['val']


trainloader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=32)
testloader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=32)
validationloader = torch.utils.data.DataLoader(dataset=validation_dataset, batch_size=32)
```

Visualising Training data


```
In [ ]: random_generated_number = torch.LongTensor(1).random_(0,(train_dataset.data
plt.imshow((trainloader.dataset[random_generated_number.numpy()[0]][0]).r
print(train_dataset.classes[trainloader.dataset[random_generated_number.n
plt.show()
random_generated_number = torch.LongTensor(1).random_(0,(train_dataset.da
plt.imshow((trainloader.dataset[random_generated_number.numpy()[0]][0]).r
print(train_dataset.classes[trainloader.dataset[random_generated_number.n

plt.show()
random_generated_number = torch.LongTensor(1).random_(0,(train_dataset.da
plt.imshow((trainloader.dataset[random_generated_number.numpy()[0]][0]).r
print(train_dataset.classes[trainloader.dataset[random_generated_number.n
```


Strawberry



Banana



Banana



We will use gray scale "1" channel as it will help us speed up process instead of doing it on "3" channels. But during visualization we will use 3 channels for color image display

Defining our neural network model

```
In [ ]: #temporary variables to store data and labels
data = []
labels = torch.FloatTensor()
inputs = torch.FloatTensor()
valid_loss = []
train_loss = []
```

```
In [ ]: def make_trainloader(train_dataset, test_dataset, validation_dataset, train_loader, test_loader, validation_loader):
    """This function will create iterable dataloaders for dataset
    Args:
        Self explanatory
    Return:
        training set data loader , validation set dataloader and test loader
    """
    trainloader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=1, shuffle=True)
    testloader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=1, shuffle=True)
    validationloader = torch.utils.data.DataLoader(dataset=validation_dataset, batch_size=1, shuffle=True)

    return trainloader, testloader, validationloader
```

We will push variables to gpu if available using `.cuda()`

Only the model and variables will be pushed to gpu

```

In [ ]: def forward_backprop(cost_fn,optimizer,model,inputs,labels,learning_rate)
        """This funtion will perform forward pass and backward pass
        Args:
            model: neural network model
            inputs : dataset
            labels : targets
            laerning_rate : hyperparameter
        Return:
            model : trained model (Single iteration), this model w
            temp_out : output from forward pass
            loss : loss calculated from loss funtion

        """

        #Taking out a temprary output from model's forward pass
        temp_out = model(inputs)
        #Resetting the current gradients for new gradients calculation
        optimizer.zero_grad()

        loss = cost_fn(temp_out , labels.type(torch.LongTensor))
        #back prop
        loss.backward()

        #Updating parameters
        optimizer.step()
        return model,temp_out,loss

```

```

In [ ]: def training_accuracy(temp_out,labels,loss):
        """Computes training accuracy and loss"""
        total1=0
        correct1=0
        _,pred = torch.max(temp_out,1)
        total1 += labels.size(0)
        if torch.cuda.is_available():
            correct1+= ((pred.type(torch.FloatTensor)).cpu() == labels.cpu())

        else:
            correct1+= (pred.type(torch.FloatTensor) == labels).sum()
            #Average accuracy
            acc1 = 100 * correct1/total1
            print('Training loss is ',loss.data[0],'Training accuracy is ',acc1)

def validating(model,validationloader,cost_fn):
    """Computes accuracy and loss of validation set"""
    correct = 0
    total = 0
    for batch_idx1 , dat1 in enumerate(validationloader):
        data1 = dat1
        inputs1 = data1[0].type(torch.FloatTensor)
        labels1 = data1[1].type(torch.FloatTensor)
        if torch.cuda.is_available():
            inputs1 = Variable(inputs1.cuda())
        else:
            inputs1 = Variable(inputs1)

        temp_out1 = model(inputs1)
        loss1 = cost_fn(temp_out1 , labels1.type(torch.LongTensor))
        valid_loss.append(loss1.data[0])
        _,pred = torch.max(temp_out1,1)
        total += labels1.size(0)
        if torch.cuda.is_available():
            correct+= ((pred.type(torch.FloatTensor)).cpu() == labels1.cpu())

        else:
            correct+= (pred.type(torch.FloatTensor) == labels1).sum()

    acc = 100 * correct/total
    print('Loss on validation set ',loss1.data[0],' Accuracy ',acc[0],'%')

```

initialize model with hyper parameters

in_channel = 1 as image was converted to gray scale

kernel are for filter size

outchannels are for total number for kernels stack

Stride has been set to 1

```

In [ ]: model = FNet(in_channel=1,out_channels1=16,out_channels2=32,kernel1=5,pool1=2)
        if torch.cuda.is_available():
            model.cuda()

```

```
In [ ]: n_epochs = 3
cost_fn = nn.CrossEntropyLoss()
learning_rate = 0.01
```

We used Adam optimizer of SGD, as it will handle out the both exploding and vanishing gradient, the algorithm is designed from ada grad algorithm but optimised for non convex funtion as the baies corrections are introduced we should not worry about saddle gradients. We set the weight decay for L2 norm

```
In [ ]: optimizer = torch.optim.Adam(model.parameters() , lr = learning_rate,we
```

```
In [ ]: #Parameter to decide wheather to print training loss on every iteration n
debug = False
#Parameter to decide wheather to save model after training or not
save = False
```

Training our models

```
In [ ]: for epochs in range(n_epochs):
    for batch_idx , dat in enumerate(trainloader):
        data = dat
        inputs = data[0].type(torch.FloatTensor)
        labels = data[1].type(torch.FloatTensor)
        if torch.cuda.is_available():
            inputs = Variable(inputs.cuda())
            labels = Variable(labels.cuda())
        else:
            inputs = Variable(inputs)
            labels = Variable(labels)

        model,temp_out,loss = forward_backprop(cost_fn,optimizer,model,inputs,labels)
        train_loss.append(loss,data[0])
        if (debug==True):
            training_accuracy(temp_out,labels,loss)

    #This validation set will be tested after each epoch
    validating(model,validationloader,cost_fn)
```

```
Loss on valdation set  tensor(0.0000)   Accuracy  tensor(97) %
Loss on valdation set  tensor(0.0000)   Accuracy  tensor(99) %
Loss on valdation set  tensor(9.5367e-07) Accuracy  tensor(99) %
```

```
In [ ]: def test_mode(data,model):
        """Testing model on new data that the model have not seen before
        Args:
            data: test data
            model : trained model
        """
        correct = 0
        total = 0
        for batch_idx1 , dat1 in enumerate(data):
            data1 = dat1
            inputs1 = data1[0].type(torch.FloatTensor)
            labels1 = data1[1].type(torch.FloatTensor)
            #GPU CHECK
            if torch.cuda.is_available():
                inputs1 = Variable(inputs1.cuda())
            else:
                inputs1 = Variable(inputs1)

            temp_out1 = model(inputs1)

            _,pred = torch.max(temp_out1,1)
            total += labels1.size(0)
            if torch.cuda.is_available():
                correct+= ((pred.type(torch.FloatTensor)).cpu() == labels1.type(torch.FloatTensor)).sum()
            else:
                correct+= (pred.type(torch.FloatTensor) == labels1.type(torch.FloatTensor)).sum()

        #AVERAGE ACCURACY
        acc = 100 * correct/total
        print('Accuracy of model on new unseen data is ',acc,'%')
```

```
In [ ]: from sklearn.metrics import confusion_matrix,classification_report

def test_confusion_matrix(data,model):
    """ this funtion will make prediction on test dataset and return pre
    Args:
        data: test data
        model : trained model
    """
    y_pred = []
    y_true = []
    for batch_idx1 , dat1 in enumerate(data):
        data1 = dat1
        inputs1 = data1[0].type(torch.FloatTensor)
        labels1 = data1[1].type(torch.FloatTensor)
        #GPU CHECK
        if torch.cuda.is_available():
            inputs1 = Variable(inputs1.cuda())
        else:
            inputs1 = Variable(inputs1)

        temp_out1 = model(inputs1)

        _,pred = torch.max(temp_out1,1)
        y_pred.append(pred)
        y_true.append(labels1)

    return y_pred,y_true

y_hat , y = test_confusion_matrix(testloader,model)
```




```
In [ ]: print( confusion_matrix( y_true = y,y_pred = y_hat) ,end = '\n\n\n')
print( classification_report( torch.tensor(y),torch.tensor(y_hat) ))
```

```
[[106   0   0   0   0]
 [  4  82   0   0   0]
 [  0   0  95   0   0]
 [  0   0   0  98   0]
 [  0   0   0   0 104]]
```

	precision	recall	f1-score	support
0.0	0.96	1.00	0.98	106
1.0	1.00	0.95	0.98	86
2.0	1.00	1.00	1.00	95
3.0	1.00	1.00	1.00	98
4.0	1.00	1.00	1.00	104
avg / total	0.99	0.99	0.99	489

```
In [ ]: sns.scatterplot(y,y_hat,color='red',markers='X')
plt.plot(y,y)
```

```
Out[ ]: [<matplotlib.lines.Line2D at 0x7fd398c57fd0>]
```




From the confusion matrix we can see that there was only 4 mislabeled datas in test dataset out of 489 , that was actually one label which should be of category 1 but was marked as category 0 , thats a 0.817% error rate , rather than that we have a great classifier model.

Note: this conclusion was generated during our run , values will change when the model is trained again.

Visualising Losses

```
In [ ]: plt.plot(range(len(train_loss)) , train_loss )
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.title("Training loss")
```


```
Out[ ]: Text(0.5,1,'Training loss')
```



Spikes in training loss is decreasing that tell us that our model was fitted well to training data

```
In [ ]: plt.plot(range(len(valid_loss)) , valid_loss )
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.title("Validation loss")
```

```
Out[ ]: Text(0.5,1,'Validation loss')
```



Spikes in validation loss is decreasing that tell us that our model was not overfitted to training data.

```
In [ ]: #Saving whole model
if(save == True):
    torch.save(model,f = 'model_beta.pt')
#Testing model on new test data
test_mode(testloader,model)
```

Accuracy of model on new unseen data is tensor(99) %

Visualizing test set on model

Now we will pick seven random images from test set as it is not seen by model. We will first plot the image, then display corresponding class. Then we will pass the pixels of the image we got from random selection in test set to our model, our model will output a vector of probability and then we will take the index of maximum probability and use that as model's prediction to that image.

```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data

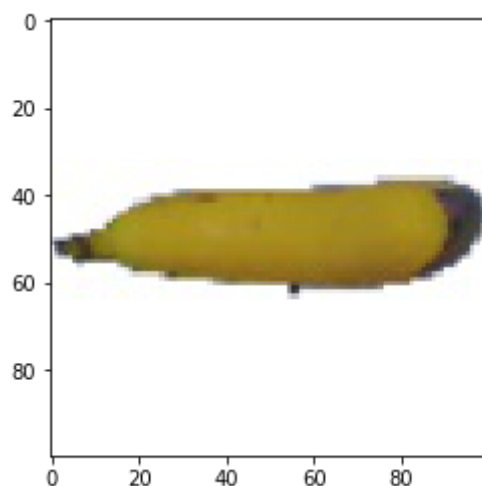
"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Banana
Actual result: Banana



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data


"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Strawberry
Actual result: Strawberry



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data

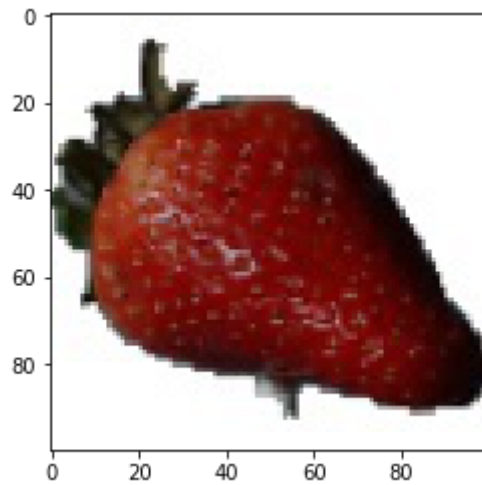
"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Strawberry
Actual result: Strawberry



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data


"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Pineapple
Actual result: Pineapple



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data


"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Pineapple
Actual result: Pineapple



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data

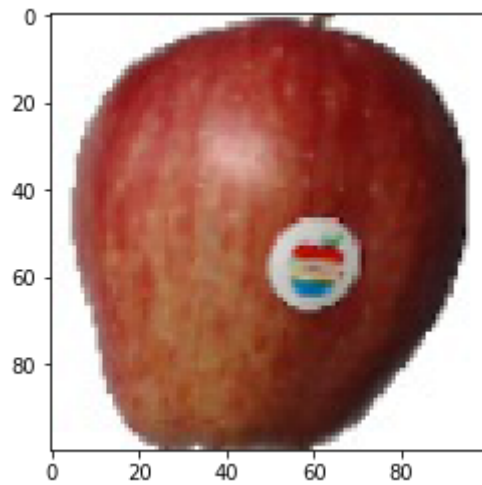
"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Apple
Actual result: Apple



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data

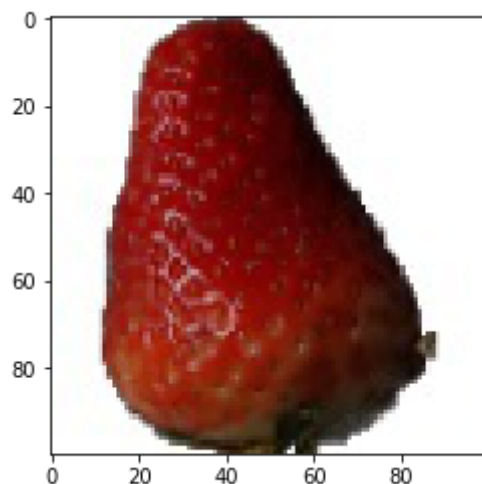
"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Strawberry
Actual result: Strawberry



```
In [ ]: """Random number generation"""
random_generated_number = torch.LongTensor(1).random_(0,(test_dataset.data


"""We use the random number as index in test set and display corresponding
temp_path = test_dataset.data['path'][random_generated_number]
img = Image.open(temp_path)
plt.imshow(img)

"""We extract the pixel from corresponding image and pass it to model for
temp_pixel = testloader.dataset[random_generated_number.numpy()[0]][0]
temp_out = model(temp_pixel.reshape(1,1,100,100))

"""We convert probability to a single prediction, where the prediction will
idx = torch.argmax(temp_out)

"""Displaying result"""
print('Expected :',(test_dataset.classes[testloader.dataset[random_genera
print('Actual result: ',test_dataset.classes[idx])
```

Expected : Pineapple
Actual result: Pineapple



Conclusion